

XML Indexing and Search

Native XML Database

- Uses XML document as logical unit
- Should support
 - Elements
 - Attributes
- Contrast with
 - DB modified for XML
 - Generic IR system modified for XML


Semi structured Retrieval

XML documents are semi structured.

Can be viewed as an ordered, labeled tree:

- *leaf nodes* containing text
- *internal nodes* define the roles of the leaf nodes in the document.
- the labeled internal nodes encode either the structure of the document (*title*, *act*, and *scene*) or metadata functions (*author*).

```
<play>
    <author>Shakespeare</author>
    <title>Macbeth</title>
    <act number="I">
        <scene number="VII">
            <title>Macbeth's castle</title>
            <verse>Will I with wine and wassail ...</verse>
        </scene>
    </act>
</play>
```



- The nodes of the tree are *XML elements* and are written with an opening and closing *tag*.
- An element can have one or more *XML attributes*.
- XML documents are processed and accessed through XML Document Object Model or *DOM*.
- DOM represents elements, attributes and text within elements as nodes in a tree.

- *XPath* is the standard for paths in XML (referred as Context)
- Ex:
act/scene -selects all *scene* elements whose parent is an *act* element.
- play//scene selects all *scene* elements occurring in a *play* element.
- /play/title
- /play//title
- /scene/title selects no elements.

XML Indexing and Search

- Most native XML databases have taken a DB approach
 - Exact match
 - Evaluate path expressions
 - No IR type relevance ranking
- Only a few that focus on relevance ranking

Data vs. Text-centric XML

- Data-centric XML: used for messaging between enterprise applications
 - Mainly a recasting of relational data
- Content-centric XML: used for annotating content
 - Rich in text
 - Demands good integration of text retrieval functionality
 - E.g., find me the ISBN #s of Books with at least three Chapters discussing cocoa production, ranked by Price

IR XML Challenge 1: Term Statistics

- There is no document unit in XML
- How do we compute tf and idf?
- Global tf/idf over all text context is useless
- Indexing granularity

IR XML Challenge 2: Fragments

- IR systems don't store content (only index)
- Need to go to document for retrieving/
displaying fragment
 - E.g., give me the Abstracts of Papers on
existentialism
 - Where do you retrieve the Abstract from?
- Easier in DB framework

IR XML Challenges 3: Schemas

- Ideally:
 - There is one schema
 - User understands schema
- In practice: rare
 - Many schemas
 - Schemas not known in advance
 - Schemas change
 - Users don't understand schemas
- Need to identify similar elements in different schemas
 - Example: employee

IR XML Challenges 4: UI

- Help user find relevant nodes in schema
 - Author, editor, contributor, “from:”/sender
- What is the query language you expose to the user?
 - Specific XML query language? No.
 - Forms? Parametric search?
 - A textbox?
- In general: design layer between XML and user

IR XML Challenges 5: using a DB

- Why you don't want to use a DB
 - Spelling correction
 - Mid-word wildcards
 - Contains vs “is about”
 - DB has no notion of ordering
 - Relevance ranking

XIRQL




XIRQL

- University of Dortmund
 - Goal: open source XML search engine
- Motivation
 - “Returnable” fragments are special
 - E.g., don’t return a <bold> some text </bold> fragment
 - Structured Document Retrieval Principle
 - Empower users who don’t know the schema
 - Enable search for any person no matter how schema encodes the data
 - Don’t worry about attribute/element

Atomic Units

- Specified in schema
- Only atomic units can be returned as result of search (unless unit specified)
- Tf.idf weighting is applied to atomic units

XIRQL Indexing




Structured Document Retrieval Principle

- A system should always retrieve the most specific part of a document answering a query.
- Example query: xql
- Document:

```
<chapter> 0.3 XQL
<section> 0.5 example </section>
<section> 0.8 XQL 0.7 syntax </section>
</chapter>
```
- Return section, not chapter

-
- motivates a retrieval strategy that returns the smallest unit that contains the information
 - hard to implement this principle algorithmically.



Text-Centric XML Retrieval

Text-centric XML retrieval


- Documents marked up as XML
 - E.g., assembly manuals, journal issues ...
- Queries are user information needs
 - E.g., give me the `Section` (`element`) of the document that tells me how to change a brake light
- Different from well-structured XML queries where you tightly specify what you're looking for.

Vector spaces and XML


- Vector spaces – tried+tested framework for keyword retrieval
 - Other “bag of words” applications in text: classification, clustering ...
- For text-centric XML retrieval, can we make use of vector space ideas?
- Challenge: capture the structure of an XML document in the vector space.

Vector spaces and XML

- For instance, distinguish between the following two cases



Content-rich XML: representation





Encoding the Gates differently

- What are the axes of the vector space?
- In text retrieval, there would be a single axis for **Gates**
- Here we must separate out the two occurrences, under Author and Title
- Thus, axes must represent not only terms, but something about their position in an XML tree


Queries

- Before addressing this, let us consider the kinds of queries we want to handle



Query types


- The preceding examples can be viewed as sub-trees of the document
- But what about?



- (**Gates** somewhere underneath Book)

Subtrees and structure


- Consider all subtrees of the document that include at least one lexicon term:



Structural terms


- Call each of the resulting subtrees a *structural term*
- Note that structural terms might occur multiple times in a document
- Create one axis in the vector space for each distinct structural term
- Weights based on frequencies for number of occurrences (just as we had tf)
- All the usual issues with terms (stemming? Case folding?) remain

Example of *tf* weighting



- Here the structural terms containing **to** or **be** would have more weight than those that don't

Down-weighting



- For the doc on the left: in a structural term rooted at the node *Play*, shouldn't ***Hamlet*** have a higher *tf* weight than ***Yorick***?
- Idea: multiply *tf* contribution of a term to a node k levels up by γ^k , for some $\gamma < 1$.

Down-weighting example, $\gamma=0.8$

- For the doc on the previous slide, the *tf* of
 - **Hamlet** is multiplied by 0.8
 - **Yorick** is multiplied by 0.64
- in any structural term rooted at Play.



The number of structural terms

- Can be huge!
- Impractical to build a vector space index with so many dimensions

Structural terms: docs+queries

- The notion of structural terms is independent of any schema/DTD for the XML documents
- Well-suited to a heterogeneous collection of XML documents
- Each document becomes a vector in the space of structural terms
- A query tree can likewise be factored into structural terms
 - And represented as a vector
 - Allows weighting portions of the query

Example query



Weight propagation


- The assignment of the weights 0.6 and 0.4 in the previous example to subtrees was simplistic
 - Can be more sophisticated
 - Think of it as generated by an application, not necessarily an end-user
- Queries, documents become normalized vectors
- Retrieval score computation “just” a matter of cosine similarity computation

Restrict structural terms?

- Depending on the application, we may restrict the structural terms
- E.g., may never want to return a Title node, only Book or Play nodes
- So don't enumerate/index/retrieve/score structural terms rooted at some nodes

The catch remains

- This is all very promising, but ...
- How big is this vector space?
- Can be exponentially large in the size of the document
- Cannot hope to build such an index
- And in any case, still fails to answer queries like



Two solutions


- Query-time materialization of axes
- Restrict the kinds of subtrees to a manageable set

Query-time materialization

- Instead of enumerating all structural terms of all docs (and the query), enumerate only for the query
 - The latter is hopefully a small set
- Now, we're reduced to checking which structural term(s) from the query match a subtree of any document
- This is tree pattern matching: given a *text tree* and a *pattern tree*, find matches
 - Except we have many text trees
 - Our trees are labeled and weighted


Example

Text =




- Here we seek a doc with **Hamlet** in the title
- On finding the match we compute the cosine similarity score
- After all matches are found, rank by sorting

Query =



(Still infeasible)

- A doc with ***Yorick*** somewhere in it:
- Query =




- Will get to it ...

Restricting the subtrees


- Enumerating all structural terms (subtrees) is prohibitive, for indexing
 - Most subtrees may never be used in processing any query
- Can we get away with indexing a restricted class of subtrees
 - Ideally – focus on subtrees likely to arise in queries

JuruXML (IBM Haifa)



- Only paths including a lexicon term
- In this example there are only 14 (why?) such paths
- Thus we have 14 structural terms in the index


Variations



- Could have used other subtrees – e.g., all subtrees with two siblings under a node
- Which subtrees get used: depends on the likely queries in the application
- Could be specified at index time – area with little research so far

Descendants


- Return to the descendant examples:



No known DTD.
Query seeks **Gates** under Author.

Handling descendants in the vector space

- Devise a *match* function that yields a score in [0,1] between structural terms
- E.g., when the structural terms are paths, measure overlap




How do we use this in retrieval?


- First enumerate structural terms in the query
- Measure each for *match* against the dictionary of structural terms
 - Just like a postings lookup, except not Boolean (does the term exist)
 - Instead, produce a score that says “80% close to this structural term”, etc.
- Then, retrieve docs with that structural term, compute cosine similarities, etc.



query



doc1



doc2

Figure $\text{cr}(q, d1) = 3/4 = 0.75$ and $\text{cr}(q, d2) = 3/5 = 0.6$.

$$\text{cr}(q, d) = \frac{1 + |q|}{1 + |d|}$$


context resemblance similarity

$$\text{sim-cr}(q, d) = \frac{\sum_{t \in V} \sum_{c_k \in C} \sum_{c_l \in C} \text{weight}(q, t, c_k) \text{weight}(d, t, c_l) \text{cr}(c_k, c_l)}{|\vec{q}| |\vec{d}|}$$

V is the vocabulary of (non-structural) terms;

C is the set of all contexts (paths) occurring in the collection and the queries; and $\text{weight}(d, t, ck)$ is the weight of term t in context ck in document d .

Example of a retrieval step



Now rank the Doc's by cosine similarity;
e.g., Doc9 scores 0.578.

- ST in the query occurs as ST1 in the index and has a match coefficient of 0.6 with a second term ST5 in the index.
- highest ranking document is Doc9 with a similarity of $1.0 \times 0.2 + 0.6 \times 0.6 = 0.56$.
- Query weights are assumed to be 1.0.

Indexing and search in JuruXML.

Indexing

For each indexing unit i :

 Compute structural terms for i

 Construct index

Search

Compute structural terms for query

For each structural term t :

 Find matching structural terms in dictionary

For each matching structural term t' :


 Compute matching weight $\text{cr}(t, t')$

Search inverted index with computed terms and weights

Return ranked list

Closing technicalities

- But what exactly is a Doc?
- In a sense, an entire corpus can be viewed as an XML document



What are the Doc's in the index?

- Anything we are prepared to return as an answer
- Could be nodes, some of their children ...


What are queries we can't handle using vector spaces?

- Find figures that describe the Corba architecture and the paragraphs that refer to those figures
 - Requires JOIN between 2 tables
- Retrieve the titles of articles published in the Special Feature section of the journal *IEEE Micro*
 - Depends on order of sibling nodes.

Can we do IDF?

- Yes, but doesn't make sense to do it corpus-wide
- Can do it, for instance, within all text under a certain element name say Chapter
- Yields a *tf-idf* weight for each lexicon term under an element
- Issues: how do we propagate contributions to higher level nodes.

Example



- Say **Gates** has high IDF under the Author element
- How should it be *tf-idf* weighted for the Book element?
- Should we use the *idf* for **Gates** in Author or that in Book?


INEX: a benchmark for text-centric XML retrieval



INEX(*IN*itiative for the *Evaluation of XML retrieval*)

- Benchmark for the evaluation of XML retrieval
 - Analog of TREC (recall CS276A)
- Consists of:
 - Set of XML documents
 - Collection of retrieval tasks

Schema of INEX Document



Fm – front matter, ti – title, atl- article title,

- There are two types of queries, called *topics*, in INEX:
 - content-only or CO topics and
 - content-and-structure or CAS topics.

Content only queries


Standard IR queries but we retrieve document component

Content and Structure (CAS) Queries


Put on constraints on which type of components are retrieved

e.g. “sections of an article in ... about”

Conceptual model for XML retrieval



Example of XML approaches



? = Aggregated weight of t_i in Document based on the instances of t_i in the sub-elements (Title, Section_1 and Section_2)

INEX


- Each engine indexes docs
- Engine team converts retrieval tasks into queries
 - In XML query language understood by engine
- In response, the engine retrieves not docs, but elements within docs
 - Engine ranks retrieved elements

INEX assessment

- For each query, each retrieved element is human-assessed on two measures:
 - Topical Relevance – how relevant is the retrieved element
 - Component Coverage – is the retrieved element too specific, too general, or just right
 - E.g., if the query seeks a definition of the Fast Fourier Transform, do I get the equation (too specific), the chapter containing the definition (too general) or the definition itself
- These assessments are turned into composite precision/recall measures

- The component coverage dimension evaluates whether the element retrieved is “structurally” correct, i.e., neither too low nor too high in the tree.
- The topical relevance dimension also has four levels: highly relevant (3) fairly relevant (2) marginally relevant (1) and irrelevant (0).

2D relevance scale



INEX corpus

- 12,107 articles from IEEE Computer Society publications
- 494 Megabytes
- Average article: 1,532 XML nodes
 - Average node depth = 6.9

INEX topics

- Each topic is an information need, one of two kinds:
 - Content Only (CO) – free text queries
 - Content and Structure (CAS) – explicit structural constraints, e.g., containment conditions.

Sample INEX CO topic

<**Title**> computational biology </**Title**>

<**Keywords**> computational biology, bioinformatics, genome, genomics, proteomics, sequencing, protein folding </**Keywords**>

<**Description**> Challenges that arise, and approaches being explored, in the interdisciplinary field of computational biology</**Description**>

<**Narrative**> To be relevant, a document/component must either talk in general terms about the opportunities at the intersection of computer science and biology, or describe a particular problem and the ways it is being attacked. </**Narrative**>

INEX assessment

- Each engine formulates the topic as a query
 - E.g., use the keywords listed in the topic.
- Engine retrieves one or more elements and ranks them.
- Human evaluators assign to each retrieved element relevance and coverage scores.

Assessments

- Relevance assessed on a scale from Irrelevant (scoring 0) to Highly Relevant (scoring 3)
- Coverage assessed on a scale with four levels:
 - No Coverage (N: the query topic does not match anything in the element)
 - Too Large (The topic is only a minor theme of the element retrieved)
 - Too Small (S: the element is too small to provide the information required)
 - Exact (E).
- So every element returned by each engine has ratings from $\{0,1,2,3\} \times \{N,S,L,E\}$

Combining the assessments

- Define scores:

$$f_{strict}(rel, cov) = \begin{cases} 1 & \text{if } rel, cov = 3E \\ 0 & \text{otherwise} \end{cases}$$

$$f_{generalized}(rel, cov) = \begin{cases} 1.00 & \text{if } rel, cov = 3E \\ 0.75 & \text{if } rel, cov \in \{2E, 3L, 3S\} \\ 0.50 & \text{if } rel, cov \in \{1E, 2L, 2S\} \\ 0.25 & \text{if } rel, cov \in \{1S, 1L\} \\ 0.00 & \text{if } rel, cov = 0N. \end{cases}$$

The f -values

- Scalar measure of goodness of a retrieved elements
- Can compute f -values for varying numbers of retrieved elements 10, 20 ... etc.
 - Means for comparing engines.

From raw f -values to ... ?

- INEX provides a method for turning these into precision-recall curves
- “Standard” issue: only elements returned by some participant engine are assessed